
Assignment Four

Matthew C. Scicluna
Département d'Informatique et de Recherche Opérationnelle
Université de Montréal
Montréal, QC H3T 1J4
matthew.scicluna@umontreal.ca

December 14, 2017

1 Entropy and Mutual Information

Let X be a discrete random variable on a finite space \mathcal{X} with $|\mathcal{X}| = k$.

1. (a) Prove that the entropy $H(X) \geq 0$, with equality only when X is a constant.

PROOF: WLOG we can assume that $p(x) > 0 \forall x \in \mathcal{X}$, (using that $0 \cdot \log 0 = 0$). We have that $H(X) = -\sum_x p(x) \log p(x) = \sum_x p(x) \log p(x)^{-1} \geq 0$, since $p(x) > 0$ and $p(x)^{-1} \geq 1$. If $H(X) = 0$ then $\exists \alpha$ such that $p(\alpha)^{-1} = 1 \Rightarrow p(\alpha) = 1$. Hence X must be a constant, as needed.

- (b) Let $X \sim p$ and q be the Uniform distribution on \mathcal{X} . What is the relation between $D(p||q)$ and $H(X)$.

CLAIM: $D(p||q) = -H(X) + \log k$

PROOF:

$$\begin{aligned} D(p||q) &= \sum_x p(x) \log \frac{p(x)}{q(x)} \\ &= \sum_x p(x) \log p(x) - \sum_x p(x) \log q(x) \\ &= -H(X) + \sum_x p(x) \log k \\ &= -H(x) + \log k \end{aligned}$$

- (c) An upper bound for $H(X)$ is $\log k$ since $H(X) = \log k - D(p||q) \Rightarrow H(X) \leq \log k$.

We consider a pair of discrete random variables (X_1, X_2) defined over the finite set $\mathcal{X}_1 \times \mathcal{X}_2$. Let $p_{1,2}$, p_1 and p_2 denote respectively the joint distribution, the marginal distribution of X_1 and the marginal distribution of X_2 . Define the mutual information as:

$$I(X_1, X_2) = \sum_{(x_1, x_2) \in \mathcal{X}_1 \times \mathcal{X}_2} p_{1,2}(x_1, x_2) \log \frac{p_{1,2}(x_1, x_2)}{p_1(x_1)p_2(x_2)}$$

We again assume WLOG that $p(x_1, x_2) > 0 \forall (x_1, x_2) \in \mathcal{X}_1 \times \mathcal{X}_2$

2. (a) CLAIM: $I(X_1, X_2) \geq 0$

PROOF: Notice that $I(X_1, X_2) = D(p_{1,2}||p_1p_2) \geq 0$ by the positiveness of $D(\cdot||\cdot)$.

- (b) We want to express $I(X_1, X_2)$ as a function of $H(X_1)$, $H(X_2)$ and $H(X_1, X_2)$.

$$\begin{aligned} I(X_1, X_2) &= \sum_{(x_1, x_2) \in \mathcal{X}_1 \times \mathcal{X}_2} p_{1,2}(x_1, x_2) \log \frac{p_{1,2}(x_1, x_2)}{p_1(x_1)p_2(x_2)} \\ &= \sum_{(x_1, x_2) \in \mathcal{X}_1 \times \mathcal{X}_2} p_{1,2}(x_1, x_2) \log p_{1,2}(x_1, x_2) - p_{1,2}(x_1, x_2) \log p_1(x_1)p_2(x_2) \\ &= -H(X_1, X_2) - \sum_{(x_1, x_2) \in \mathcal{X}_1 \times \mathcal{X}_2} \left(p_1(x_1)p_2(x_2) \log p_1(x_1) - p_1(x_1)p_2(x_2) \log p_2(x_2) \right) \\ &= -H(X_1, X_2) - \sum_{j=1}^2 \sum_{(x_1, x_2) \in \mathcal{X}_1 \times \mathcal{X}_2} p_1(x_1)p_2(x_2) \log p_j(x_j) \\ &= -H(X_1, X_2) - \sum_{j=1}^2 \sum_{x_j \in \mathcal{X}_j} p_j(x_j) \log p_j(x_j) \\ &= -H(X_1, X_2) + H(X_1) + H(X_2) \end{aligned}$$

And so we can represent $I(X_1, X_2)$ using $H(X_1)$, $H(X_2)$ and $H(X_1, X_2)$, as needed.

- (c) From the previous result we have that $I(X_1, X_2) \geq 0 \Rightarrow H(X_1) + H(X_2) \geq H(X_1, X_2)$, and so the maximal entropy of (X_1, X_2) is $H(X_1) + H(X_2)$. By definition this only occurs when $I(X_1, X_2) = 0$, which only occurs if $p_{1,2}(x_1, x_2) = p_1(x_1)p_2(x_2) \forall x_1, x_2 \in \mathcal{X}_1 \times \mathcal{X}_2$. This can be seen directly from the definition of I and using the strict positivity of $p(x_1, x_2)$.

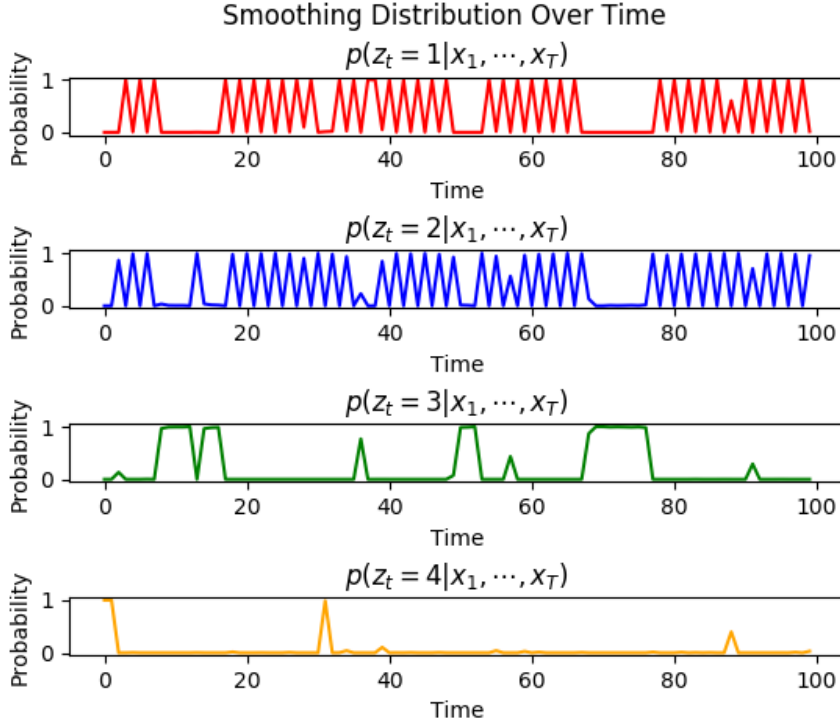


Figure 2.1: The smoothing distribution for the first 100 time points, using the parameters learned in the previous homework.

2 HMM – Implementation

We use an HMM model to account for the possible temporal structure of some data. We consider the following HMM model: the chain $(z_t)_{t=1}^T$ has $K = 4$ possible states, with an initial probability distribution $\pi \in \Delta_4$ and a probability transition matrix $A \in \mathbb{R}^{4 \times 4}$ where $A_{ij} = p(z_t = i | z_{t-1} = j)$ and conditionally on the current state z_t , we have observations obtained from Gaussian emission probabilities $x_t | (z_t = k) \sim N(x_t | \mu_k, \Sigma_k)$.

2.1 Fake Parameters Inference

We computed the vectors $\alpha(z_t) = p(z_t, x_{1:t})$ and $\beta(z_t) = p(x_{(t+1):T} | z_t)$ on the test set from Assignment 3 using the following parameters: $\pi_k = \frac{1}{4}$, $A_{ii} = \frac{1}{2}$ and $A_{ij} = \frac{1}{6}$, $i \neq j$, and μ_k, Σ_k as defined in the homework. We used these to compute the posterior of the latent variable over time $p(z_t | x_1, \dots, x_T)$. We plotted the first 100 datapoints in figure 2.1.

2.2 M-Step Derivation

We now derive the M-Step for the Hidden Markov Model.

Let $\theta^{(s)} = (\pi^{(s)}, A^{(s)}, \mu_1^{(s)}, \dots, \mu_K^{(s)}, \Sigma_1^{(s)}, \dots, \Sigma_K^{(s)})$ be the ML parameters learned during step s of EM. Let $\gamma_{tk} = P(z_t = k | x_{1:T})$ and $\xi_{tlm} = P(z_t = l, z_{t+1} = m | x_{1:T})$ - which are the quantities that were computed in the E-Step using $\theta^{(s)}$. The Expected Complete Data Log-Likelihood (at step $s + 1$) is:

$$Q(\theta, \theta^{(s)}) = \sum_{k=1}^K \gamma_{1k} \log \pi_k + \sum_{t=1}^T \sum_{k=1}^K \gamma_{tk} \log P(\bar{x}_t | \mu_k, \Sigma_k) + \sum_{t=1}^{T-1} \sum_{l=1}^K \sum_{m=1}^K \xi_{tlm} \log A_{lm} \quad (2.1)$$

To solve for π_k we look for stationary points, subject to the constraint $\sum_{k=1}^K \pi_k = 1$

$$\frac{\partial}{\partial \pi_k} Q(\theta, \theta^{(s)}) - \lambda \left(\sum_{k=1}^K \pi_k - 1 \right) = \frac{\gamma_{1k}}{\pi_k} - \lambda = 0 \quad (2.2)$$

$$\Rightarrow \pi_k = \frac{\gamma_{1k}}{\lambda} \quad (2.3)$$

And using that $\sum_{k=1}^K \pi_k = 1$ we have that:

$$\pi_k^{(s+1)} = \frac{\gamma_{1k}}{\sum_{l=1}^K \gamma_{1l}} \quad (2.4)$$

To solve for A_{lm} we look for stationary points, subject to the constraint $\sum_{l=1}^K A_{lm} = 1$

$$\frac{\partial}{\partial A_{lm}} Q(\theta, \theta^{(s)}) - \lambda \left(\sum_{l=1}^K A_{lm} - 1 \right) = \sum_{t=1}^{T-1} \frac{\xi_{tlm}}{A_{lm}} - \lambda = 0 \quad (2.5)$$

$$\Rightarrow A_{lm} = \frac{\sum_{t=1}^{T-1} \xi_{tlm}}{\lambda} \quad (2.6)$$

And using that $\sum_{l=1}^K A_{lm} = 1$ we have that:

$$A_{lm}^{(s+1)} = \frac{\sum_{t=1}^{T-1} \xi_{tlm}}{\sum_{l=1}^K \sum_{t=1}^{T-1} \xi_{tlm}} \quad (2.7)$$

To solve for μ_k we look for stationary points

$$\frac{\partial}{\partial \mu_k} Q(\theta, \theta^{(s)}) = \sum_{t=1}^T \frac{\partial}{\partial \mu_k} \frac{-\gamma_{tk}}{2} (\bar{x}_t - \mu_k)^T \Sigma_k^{-1} (\bar{x}_t - \mu_k) \quad (2.8)$$

$$= \sum_{t=1}^T \frac{-\gamma_{tk}}{2} (\bar{x}_t - \mu_k) \Sigma_k^{-1} = 0 \quad (2.9)$$

$$\Rightarrow \sum_{t=1}^{T-1} \gamma_{tk} (\bar{x}_t - \mu_k) = 0 \quad (2.10)$$

$$\Rightarrow \mu_k^{(s+1)} = \frac{\sum_{t=1}^T \gamma_{tk} \bar{x}_t}{\sum_{t=1}^{T-1} \gamma_{tk}} \quad (2.11)$$

To solve for Σ_k we look for stationary points. As in assignment 2, we take the derivative w.r.t Σ_k^{-1}

$$\frac{\partial}{\partial \Sigma_k^{-1}} Q(\theta, \theta^{(s)}) = \sum_{t=1}^T \frac{\partial}{\partial \mu_k} \left(\frac{-\gamma_{tk}}{2} \log |\Sigma_k| - \frac{\gamma_{tk}}{2} \left(\bar{x}_t - \mu_k^{(s+1)} \right)^T \Sigma_k^{-1} \left(\bar{x}_t - \mu_k^{(s+1)} \right) \right) \quad (2.12)$$

Differentiating the first term yields

$$\frac{\partial}{\partial \Sigma_k^{-1}} \frac{-\gamma_{tk}}{2} \log |\Sigma_k| = \frac{\partial}{\partial \Sigma_k^{-1}} \frac{\gamma_{tk}}{2} \log |\Sigma_k^{-1}| \quad (2.13)$$

$$= \frac{\gamma_{tk}}{2} \Sigma_k \quad (2.14)$$

Differentiating the second term yields

$$\frac{\partial}{\partial \Sigma^{-1}} \frac{\gamma_{tk}}{2} \left(\bar{x}_t - \mu_k^{(s+1)} \right)^T \Sigma_k^{-1} \left(\bar{x}_t - \mu_k^{(s+1)} \right) \quad (2.15)$$

$$= \frac{\partial}{\partial \Sigma^{-1}} \frac{\gamma_{tk}}{2} \text{tr} \left(\left(\bar{x}_t - \mu_k^{(s+1)} \right) \left(\bar{x}_t - \mu_k^{(s+1)} \right)^T \Sigma_k^{-1} \right) \quad (2.16)$$

$$= \frac{\gamma_{tk}}{2} \left(\bar{x}_t - \mu_k^{(s+1)} \right) \left(\bar{x}_t - \mu_k^{(s+1)} \right)^T \quad (2.17)$$

Where (2.15) and (2.16) come from results proved in class and on Homework 2. After substituting (2.14) and (2.17) into (2.12) and equating to 0 we have that:

$$\sum_{t=1}^T \gamma_{tk} \Sigma_k = \sum_{t=1}^T \gamma_{tk} \left(\bar{x}_t - \mu_k^{(s+1)} \right) \left(\bar{x}_t - \mu_k^{(s+1)} \right)^T \quad (2.18)$$

$$\Rightarrow \Sigma_k^{(s+1)} = \frac{\sum_{t=1}^T \gamma_{tk} (\bar{x}_t - \mu_k^{(s+1)}) (\bar{x}_t - \mu_k^{(s+1)})^T}{\sum_{t=1}^T \gamma_{tk}} \quad (2.19)$$

2.3 Deriving Parameters using EM Algorithm

We implemented the EM algorithm to learn the parameters of the model, initializing them with the values provided in the homework. We trained the model on the EMGaussians.train dataset. The transition matrix found is as follows:

$$A = \begin{bmatrix} 0.0158 & 0.863 & 0.067 & 0.052 \\ 0.947 & 0.0226 & 0.022 & 0.0266 \\ 0.0289 & 0.027 & 0.874 & 0.0195 \\ 0.0087 & 0.0873 & 0.0371 & 0.902 \end{bmatrix}$$

The rest of the parameters are provided in table 2.1:

Table 2.1: EM Parameter Values

<i>Cluster</i>	π	μ	Σ
1	0	$\begin{bmatrix} -1.94 \\ 4.2 \end{bmatrix}$	$\begin{bmatrix} 3.34 & 0.32 \\ 0.32 & 2.84 \end{bmatrix}$
2	0	$\begin{bmatrix} 4.0 \\ 3.64 \end{bmatrix}$	$\begin{bmatrix} 0.197 & 0.275 \\ 0.275 & 12.4 \end{bmatrix}$
3	0	$\begin{bmatrix} 3.79 \\ -3.97 \end{bmatrix}$	$\begin{bmatrix} 0.95 & 0.077 \\ 0.077 & 1.58 \end{bmatrix}$
4	1	$\begin{bmatrix} -2.96 \\ -3.44 \end{bmatrix}$	$\begin{bmatrix} 6.88 & 6.66 \\ 6.66 & 6.75 \end{bmatrix}$

2.4 Log Likelihood of Train and Test Data

The log likelihood for the training and test data was computed and is displayed in figure 2.2. The log likelihood of the training set is consistently higher than on the test set. This makes sense since the model is being trained to (indirectly) maximize the likelihood of the training set.




Figure 2.2: The log likelihood computed for each iteration of the EM algorithm.

Table 2.2: Training and Test Log Likelihood for Various Models

<i>Model</i>	Train LL	Test LL
HMM	-1919	-1949
GMM (Full)	-2367	-2450
GMM (Spherical)	-2722	-2769

2.5 Comparing Previous Models

We compared the values of log-likelihoods of the Gaussian Mixture Model (with full and spherical covariance structures) and of the HMM on the training and test data. The results are presented in table 2.2.

As was expected, the training log likelihood was higher then the test log likelihood since the models were trained to maximize the training data. The HMM had the highest log likelihood, followed by the GMM with Full covariance, and the GMM with Spherical covariance. This makes sense since the data seemed to be generated using an HMM (since the data were clearly correlated through time and were clearly from different clusters). It does make sense to make this comparison because each model is a subset of the other, with the GMM with general covariance being an HMM with each column of A being identical (i.e. each z_t is independent of z_{t-1}).

2.6 Viterbi Decoding Pseudocode

Let $x_{1:T}$ be observations and $z_{1:T}$ be hidden states, where $z_t \in \{1, \dots, K\}$. Define $V, P \in \mathbb{R}^{T \times K}$ where V_{ij} is the probability of the most likely path z_1, \dots, z_i with $z_i = j$ and P_{ij} is the value of the corresponding z_{i-1} . Let $O_{ij} = P(\bar{x}_i | z_i = j)$. The algorithm is presented below.

The algorithm works since

$$\arg \max_{z_{1:T}} P(z_{1:T} | \bar{x}_{1:T}) \quad (2.20)$$

$$= \arg \max_{z_{1:T}} P(\bar{x}_{1:T}, z_{1:T}) \quad (2.21)$$

$$= \arg \max_{z_T} \left\{ \arg \max_{z_{T-1}} P(\bar{x}_T | z_T) P(z_T | z_{T-1}) \left\{ \arg \max_{z_{T-2}} P(\bar{x}_{T-1} | z_{T-1}) \cdots \right\} \right\} \quad (2.22)$$

$$= \arg \max_{z_T} \arg \max_{z_{T-1}} O_{T, z_T} A_{z_T, z_{T-1}} V_{T-1, z_{T-1}} \quad (2.23)$$

Which is essentially a forward pass with max operation instead of sum. Steps 4 – 7 computes the needed V and P . Step 8 uses these to compute (2.20), returning z_T . Steps

Algorithm 1 Calculate $\arg \max_{z_{1:T}} P(z_{1:T} | \bar{x}_{1:T})$

```

1: for state  $j \in \{1, \dots, K\}$  do
2:    $V_{1j} \leftarrow \pi_j O_{1j}$   $\triangleright P(\bar{x}_1, z_1 = j)$ 
3:    $P_{1j} \leftarrow -1$   $\triangleright$  initialize with anything
4: for time  $i \in \{2, \dots, T\}$  do
5:   for state  $j \in \{1, \dots, K\}$  do
6:      $V_{ij} \leftarrow \max_k O_{ij} A_{kj} V_{i-1,k}$ 
7:      $P_{ij} \leftarrow \arg \max_k O_{ij} A_{kj} V_{i-1,k}$   $\triangleright P(\bar{x}_i | z_i = k) P(z_i = k | z_{i-1} = j) V_{i-1,k}$ 
8:    $z_T \leftarrow \arg \max_k V_{Tk}$ 
9:   for  $i \in \{T, T-1, \dots, 2\}$  do
10:     $z_{i-1} \leftarrow V_{i, z_i}$   $\triangleright$  backtracking
return  $z_1, \dots, z_T$ 

```

9 – 10 use P and z_T to recover $z_{1:T-1}$. Specifically:

$$\begin{aligned}
z_{T-1} &= P_{T-1, z_T} \\
z_{T-2} &= P_{T-2, z_{T-1}} \\
&\dots
\end{aligned}$$

2.7 Viterbi Decoding Implementation

We implemented Viterbi decoding. We used the parameters learned with the EM algorithm to compute the most likely sequence of states with the Viterbi algorithm. The results are presented graphically in Figure 2.3.

2.8 Marginal Probability Computations

For the datapoints in `EMGaussian.test`, we computed the marginal probability $p(z_t | x_1, \dots, x_T)$ for each point to be in state $\{1, 2, 3, 4\}$ using the parameters learned on the training set. For each state, we plotted the probability of being in that state as a function of time. The results for the first 100 datapoints are presented in figure 2.4. We then computed their most likely state according to $p(z_t | x_1, \dots, x_T)$ and using Viterbi decoding. The results are presented in Figure 2.5. Note that the most likely sequence of states coincided for the first 100 datapoints (although they did differ at $t=448$).

2.9 How to Train the Model With Unknown Number of Clusters

One way to choose the number of clusters K is to treat it as a hyperparameter and select the optimum using cross validation and grid search. For each K in the hyperparameter space, a KMeans model can be trained, and (assuming a "good" minima is achieved) the cluster centers learned can be used as initial values for a GMM. The parameters learned




Figure 2.3: The training data is presented along with the cluster means, inferred from EM. The datapoints are colored based on each datapoints' Viterbi decoded cluster assignment.




Figure 2.4: The smoothing distribution for the first 100 time points of the test set, using the parameters learned with EM.




Figure 2.5: The The most likely state for the first 100 time points of the test data, using Viterbi decoding and the marginal probability.

by the GMM can be used as initial parameters for the HMM. The log likelihood can be evaluated using cross validaion and the K chosen using grid search.